

# A more Pragmatic Implementation of the Lock-free, Ordered, Linked List

Jesper Larsson Träff, Manuel Pöter  
TU Wien

Faculty of Informatics  
Institute of Computer Engineering, Research Group Parallel Computing  
Favoritenstrasse 16, 1040 Vienna, Austria

October 29th, 2020

## Abstract

The lock-free, ordered, linked list is an important, standard example of a concurrent data structure. An obvious, practical drawback of textbook implementations is that failed compare-and-swap (`CAS()`) operations lead to retraversal of the entire list (retries), which is particularly harmful for such a linear-time data structure. We alleviate this drawback by first observing that failed `CAS()` operations under some conditions do not require a full retry, and second by maintaining approximate backwards pointers that are used to find a closer starting position in the list for operation retry. Experiments with both a worst-case deterministic benchmark, and a standard, randomized, mixed-operation throughput benchmark on three shared-memory systems (Intel Xeon, AMD EPYC, SPARC-T5) show practical improvements ranging from significant, to dramatic, several orders of magnitude.

## 1 Introduction

The lock-free, ordered, singly linked list as proposed in [5, 8] is a textbook example of a concurrent data structure [6, 12]. The data structure supports lock-free insertion (`add()`) and deletion (`rem()`), and wait-free contains (`con()`) operations on items identified by a unique key. The lock-free implementation is actually quite subtle. The ordering condition and a relaxed invariant makes it possible to do with a single-word compare-and-swap operation (`CAS()`), and all operations can be shown to be linearizable even though linearization does not always happen at fixed points in the code (`CAS()` operations). The lock-free data structure has many direct and indirect applications, notably in the implementation of concurrent skiplists and hash tables [8, 11, 13, 14].

An obvious, practical drawback of the standard implementation is the draconic action on failed `CAS()` operations: A retraversal from the head of the list is simply initiated. For a linear-time data structure, this is particularly harmful since it in the worst case entails traversal of the whole list. For any individual thread, this can happen indefinitely, since the implementation is not starvation-free. Available implementations seem to implement ordered, lock-free lists in this way [1, 3, 4, 7].

In [2], this problem was addressed by maintaining a doubly linked list. The implementation is more complex involving two atomic flags and extra `CAS()` operations; the paper focuses on a

theoretical analysis, and gives no experimental results, and none seem to have been prominently reported.

The present, short paper gives some most pragmatic improvements to the textbook implementation that seem to benefit practical performance from significantly to dramatically. The improvements first seek to avoid complete retraversal of the list wherever possible by examining the reason for each failed `CAS()` operation. In both `add()`, `rem()` and the crucial internal operation for finding the position in the list on which to add or remove an item, there are such possibilities (this may well be done in other implementations, but we have not found any published claims). Second, by extending the list items with a predecessor pointer that is maintained approximately by conditional, atomic stores and loads, all failed `CAS()` operations traverse the list backwards only to a point where a new, forwards find operation can be started. The backward pointers only have to fulfill that for any list item, there is a path back to the head (sentinel) element of the list. The improvements change the instances where the list operations linearize, but we claim that the implementation remains linearizable largely as the textbook implementation.

Unfortunately, the introduction of backward pointers and cursor significantly complicate the problem of safe memory reclamation. The implementation benchmarked here does only simple memory reclamation after each experiment. How to do proper memory reclamation with these improvements is at the moment beyond this paper.

## 2 Improvement and Implementation

The standard lock-free, ordered linked list implementation maintains the invariants that items are in key order and that an item is in the list *iff* it is reachable from the head and not marked. The mark (flag) kept with each list item is a stolen bit in the item's pointer to the next larger item, such that mark and pointer can be checked and updated atomically by a single `CAS()` operation. Textbook implementations use a non-visible search function for locating the position where a given key could be; this function has the additional, crucial task of linking out items that have become marked. A `CAS()` operation is used for linking out a marked item, for inserting a new item, and for marking an item for deletion. Each of these `CAS()` operations check both that the item from which the operation is performed has not become marked, and that the pointer to the next larger item has not changed.

We observe the following.

- Upon failure of the `CAS()` operation in the search function, if the item on which the operation failed has *not* become marked (only the pointer changed due to another thread having had success in linking out the item), there is no reason to restart the search from the head of the list; it suffices to reread the next pointer of the item.
- In the `rem()` operation, if the `CAS()` operation fails due to the item having become marked, the delete can be linearized as failed, since the item has been removed by another thread. The linearization point is not the failed `CAS()`, though, but the earlier of this and the point just before an overlapping `add()` operation is about to successfully insert an item with the same key. If instead the `CAS()` failed due to the pointer to the next item having changed, the thread can simply try again to mark the item. This can be repeated until the item to be deleted has become marked, successfully by some thread. Once the node is marked it

is logically deleted and will be removed eventually. So a failure of the subsequent `CAS()` to unlink the node from the list can simply be ignored.

- In an `add()` operation, if the `CAS()` operation fails due to the next pointer having changed (but not the mark), the search for the position where to add can be continued from the item, there is no reason to go back to the head of the list. All that is needed is to reread the pointer.

The three observations lead to three mild, pragmatic improvements of the textbook implementation. The first (search function) and second (`rem()`) are illustrated with concrete code snippets from our implementation (that is available from the authors) in Listing 1 and Listing 2. The third mild improvement (in `add()`) is handled by the search function which checks the mark of the current element. Implementation is done in C11 using the standard atomic operations with the C11 memory model (acquire-release); `LOAD`, `STORE` and `CAS` are macros for these. It is worth noting that these changes do not introduce any issues regarding memory reclamation; any of the commonly used techniques will do.

The changes to the `rem()` operation makes another (platform specific) improvement possible, namely to use an atomic fetch-and-or operation to set the delete mark on the next pointer. On architectures that natively support this operation this would have the advantage that the marking operation cannot fail. We also benchmarked this potential improvement in the course of this paper by relying on the corresponding C11 operation. We note that x86 architectures can support and atomic (lock) or operation, but not an atomic fetch-and-or.

The more intrusive improvement which actually entails the milder improvements for free, extends list items with a backwards pointer to a previous, smaller key element. On a `CAS()` failure, the search function will simply have the task to go backwards in the list, that is through smaller key items, until an unmarked element is found from which the search in increasing key order can be started.

The search function with backward pointer is shown in Listing 3. In order for the implementation to be correct, the backward pointers only have to fulfill that for any item, there is a path via backward pointers back to the head of the list. When a new item is inserted into the list, the backwards pointer of the successor is set with an atomic store to point to the new element. Likewise, when an item is removed, the backwards pointer of the successor item is updated to skip the removed item. This suffices to maintain the invariant, but through long sequences of (concurrent) insertions and deletions, backwards pointers can become imprecise in that they skip many items that are actually in the list. As seen in Listing 3, we try to maintain more precise backwards pointers by updating them during forwards traversals. Since updates with atomic stores are expensive due to cache coherence activity, we only update a pointer if a test (with non-atomic, relaxed loads) shows that a pointer is not correct. Also, when a marked item is linked out, the backwards pointer of its successor is updated to skip the linked out item. One reason for this is that memory reclamation can (only) be done when there are no pointers to a removed item.

In contrast to the algorithm of [2], this optimistic implementation has no extra flag to be maintained with possibly expensive (failing) `CAS()` operations.

The final, highly pragmatic improvement is to exploit throughout the doubly linked structure of the list. Each thread maintains a cursor item in the list. On its next operation, depending on the key of the item to be located, the find operation will search either forwards (increasing key order) or backwards (smaller key order) in the list from the cursor. Each `add()`, `rem()` and `con()` operations sets the cursor to the item before the located item.

Obviously, keeping a cursor to an item in the middle of the list reduces the average runtime complexity for a singly linked structure as well, since we only have to search from head to cursor, or from cursor to end, depending on the key of the item to be located.

Unfortunately, the introduction of backward pointers significantly complicates the problem of safe memory reclamation, because it can happen that a backwards pointer references a node that has already been removed from the list. In order to retire a node for reclamation, it has to be ensured that it is not referenced by any next or backwards pointer. Proper memory reclamation with these improvements is currently outside the scope this work.

### 3 Experimental results

We have conducted a number of experiments with the described C11 implementations on different platforms to illustrate concrete, pragmatic benefits of the observations and improvements we discussed. Experiments have been done on three standard multi-core systems:

1. A 64-core AMD EPYC system with two 32-core AMD EPYC 7551 sockets at 2.9GHz
2. An 80-core Intel Xeon with 8 E7-8850 sockets at 1.06GHz
3. A 64-core SPARC v9 system with eight 8-core SPARC-T5 sockets at 3.6GHz and 8x SMT.

We have used two different benchmarks.

- Deterministic worst-case benchmark: Starting from an empty list, each thread performs the following three sequences each of length of  $n$ : `con( $k(i)$ )`, `add( $k(i)$ )`, `con( $k(i)$ )`, `add( $k(i)$ )` for  $i = 0$  to  $n - 1$ , then `con( $k(i)$ )`, `rem( $k(i)$ )`, `con( $k(i)$ )`, `rem( $k(i)$ )` from  $i = n - 1$  to  $i = 0$ , finally `con( $k(i)$ )` for  $i = 0$  to  $n - 1$ . The key function  $k(i)$  is chosen either such that each thread  $t$  has its own sequence of keys, disjoint from all other threads,  $k(i) = t + ip$  ( $p$  number of threads), or such that all threads have the same key sequences,  $k(i) = i$ . Due to the linear search in the ordered lists, the sequential behavior per thread is  $O(pn^2)$  (for disjoint keys) or  $O(n^2)$  (for same keys) steps.
- Standard random operation mix benchmark: Keys are chosen uniformly at random in an interval  $[0, U - 1]$  (`random_r`). The list is prefilled with some number of items  $f$ , then a mix of randomly chosen `add()`, `rem()` and `con()` operations is performed with predefined probability for each operation (here we report only for the mix 10 – 10 – 80). Each thread performs the same number of operations  $c$ , and threads have different seeds for the random number generator (we use the thread-safe `random_r()` `generator`, except on SPARC since this function is not available on Solaris). For chosen  $f$  and  $U$  the number of elements of the list will not vary too much.

We use OpenMP to manage threads and time the benchmarks. The benchmarks can also be configured such that each thread operates on a private list, such that there is no interaction required between threads. In this configuration, we can use either the lock-free implementation, or a standard, sequential (doubly or singly linked) list implementation. These configurations can give an idea of the system and memory overheads when there is no actual interaction between threads. We do not report on the thread private behavior here. We report on six implementation variants, namely a) *draconic* (textbook implementation), b) *singly* linked list with mild improvements, c)

Table 1: Deterministic benchmark  $k(i) = i$ , AMD EPYC system,  $p = 64, n = 100000$ . Operation breakdown: “adds” is the number of successful `add()` operations, “rems” the number of successful `rem()` operations, “cons” the number of element traversals over all `con()` operations, “trav” the number of list element traversals in the search function, “fail” the number of `CAS()` failures, and “rtry” the number of retries in the search function. Variants: a) draconic, b) singly, c) doubly, d) singly-cursor, e) singly-fetch-or, f) doubly-cursor.

Variant	Time (ms)	Total ops	Throughput (Kops/s)	adds	rems	cons	trav	fail	rtry
a)	162841.78	57600000	353.72	143909	143909	1218424056670	1218522498969	2715	2534
b)	145519.31	57600000	395.82	145091	145091	1182002909974	921060012934	4697	64
c)	141036.26	57600000	408.41	129492	129492	1271917666836	850820221201	18113	0
d)	85334.28	57600000	674.99	133575	133575	38267280	1236958075255	19376	22
e)	78687.72	57600000	732.01	123834	123834	38260212	1236723136990	16235	24
f)	481.28	57600000	119681.20	100333	100333	76141143	56187324	591036	16450

*doubly* linked list with approximate backward pointers and retry from from head of list, d) *singly-cursor* singly linked list with per thread retry from the last recorded position (cursor) in the list, e) *singly-fetch-or* singly linked list with per thread retry from the last recorded position (cursor) and atomic fetch-and-or in the `rem()` operations, and f) *doubly-cursor* doubly linked list with per thread retry from the last recorded position (cursor) in the list. On Intel and AMD we used the standard `glibc` memory allocator; on SPARC we used the `libumem` allocator. All implementations and code used in the experiments are available from the authors.

Results for fixed number of threads on the different platforms for the two benchmarks are shown in Tables 1, 2, 3, 4, 5, and 6. We report throughput over the number of operations, and also count the number of failed `CAS()` operations, the number of retries, the total number of list item traversals in the search operation, and the total number of traversals in contains operations, and the total number of successful `add()` and `rem()` operations. The results in the tables are for one single specific run of the benchmarks. Runs differ, but in a tolerable range.

We investigate the (weak) scalability of five of the six (excluding the atomic fetch-and-or variant) variants with the random operation mix benchmark. Here, we plot the mean throughput of 5 experiments. The scalability results are shown in Figures 1, 2 and 3. The experiments were run with a key range of 32768 and an update ratio of 50% (25% `add()`, 25% `rem()`); the lists were prefilled with 16384 items. These settings are comparable to those used in [3].

In the tables and plots,  $p$  denotes the number of started threads,  $n$  the list length for the deterministic benchmark,  $c$  the number of operations per thread for the random mix benchmark (thus weak scaling in the scalability experiments, since the number of operations per thread is kept fixed for increasing  $p$ ),  $f$  the number of prefilled elements, and  $U$  the upper bound for the key range.

## 4 Discussion

The mild improvements over the draconic textbook implementation clearly reduce the number of retries; as a consequence it also reduces the completion time and improves the throughput, but not as much as expected. The (emulated) atomic fetch-and-or operation as expected brings no improvement over the corresponding improved singly linked list with cursor.

Table 2: Deterministic benchmark  $k(i) = t + ip$ , AMD EPYC system,  $p = 64, n = 10000$ . Operation breakdown: “adds” is the number of successful `add()` operations, “rems” the number of successful `rem()` operations, “cons” the number of element traversals over all `con()` operations, “trav” the number of list element traversals in the search function, “fail” the number of `CAS()` failures, and “rtry” the number of retries in the search function. Variants: a) draconic, b) singly, c) doubly, d) singly-cursor, e) singly-fetch-or, f) doubly-cursor.

Variant	Time (ms)	Total ops	Throughput (Kops/s)	adds	rems	cons	trav	fail	rtry
a)	860636.43	5760000	6.69	640000	640000	808043265652	823284224246	46720	22816
b)	787290.12	5760000	7.32	640000	640000	808320969602	608143267099	38738	0
c)	807260.63	5760000	7.14	640000	640000	801938148075	602327832241	28103	0
d)	504290.83	5760000	11.42	640000	640000	201364518449	804041576627	45914	0
e)	578657.21	5760000	9.95	640000	640000	202104065277	807669700081	51486	1
f)	285.97	5760000	20142.13	640000	640000	58322194	5202640	17221	25

Table 3: Random operation mix benchmark, AMD EPYC system,  $p = 64, c = 1000000, f = 1000, U = 10000$ . Operation Mix 10% `add()`, 10% `rem()`, 80 % `con()`. Operation breakdown: “adds” is the number of successful `add()` operations, “rems” the number of successful `rem()` operations, “cons” the number of element traversals over all `con()` operations, “trav” the number of list element traversals in the search function, “fail” the number of `CAS()` failures, and “rtry” the number of retries in the search function. Variants: a) draconic, b) singly, c) doubly, d) singly-cursor, e) singly-fetch-or, f) doubly-cursor.

Variant	Time (ms)	Total ops	Throughput (Kops/s)	adds	rems	cons	trav	fail	rtry
a)	42633.85	64000000	1501.15	3198087	3203070	127961354252	32034345892	23399	23189
b)	44169.92	64000000	1448.95	3197813	3202696	127951211375	26665833669	18441	1
c)	43643.95	64000000	1466.41	3197408	3202406	128000172533	26651455978	20517	0
d)	28527.89	64000000	2243.42	3197537	3202414	85313098938	21336289626	25488	0
e)	30076.08	64000000	2127.94	3196895	3201757	85299471964	21332876996	24305	0
f)	20028.06	64000000	3195.52	3201780	3199113	85221592358	21296756185	15544	0

Table 4: Deterministic benchmark  $k(i) = i$ , Intel Xeon system,  $p = 80, n = 100000$ . Operation breakdown: “adds” is the number of successful `add()` operations, “rems” the number of successful `rem()` operations, “cons” the number of element traversals over all `con()` operations, “trav” the number of list element traversals in the search function, “fail” the number of `CAS()` failures, and “rtry” the number of retries in the search function. Variants: a) draconic, b) singly, c) doubly, d) singly-cursor, e) singly-fetch-or, f) doubly-cursor.

Variant	Time (ms)	Total ops	Throughput (Kops/s)	adds	rems	cons	trav	fail	rtry
a)	314154.80	72000000	229.19	111618	111618	1574094573004	1574104025889	3655	301
b)	264868.53	72000000	271.83	111229	111229	1553877268888	1188479836763	457	6
c)	265895.34	72000000	270.78	108673	108673	1599794312194	1151226363104	2795	1
d)	184202.03	72000000	390.88	121969	121969	47879590	1560938517400	8405	1
e)	195536.43	72000000	368.22	126514	126514	47874986	1555697504918	6189	50
f)	3366.28	72000000	21388.60	100189	100189	95827925	69781787	57781	138

Table 5: Deterministic benchmark  $k(i) = t + ip$ , Intel Xeon system,  $p = 80, n = 10000$ . Operation breakdown: “adds” is the number of successful `add()` operations, “rems” the number of successful `rem()` operations, “cons” the number of element traversals over all `con()` operations, “trav” the number of list element traversals in the search function, “fail” the number of `CAS()` failures, and “rtry” the number of retries in the search function. Variants: a) draconic, b) singly, c) doubly, d) singly-cursor, e) singly-fetch-or, f) doubly-cursor.

Variant	Time (ms)	Total ops	Throughput (Kops/s)	adds	rems	cons	trav	fail	rtry
a)	3482460.95	7200000	2.07	800000	800000	1254565652071	1446526400768	375424	114672
b)	3176148.60	7200000	2.27	800000	800000	1276788904826	957852619867	287865	1
c)	2917187.37	7200000	2.47	800000	800000	1276828251706	957639298722	254600	0
d)	2147669.87	7200000	3.35	800000	800000	319401544377	1277563045294	488450	69
e)	2129745.54	7200000	3.38	800000	800000	319037405095	1276676865674	484925	78
f)	584.10	7200000	12326.70	800000	800000	69175184	6697424	33751	439




Figure 1: Scalability with threads of random operation mix benchmark, AMD EPYC system,  $c = 50000, f = 16384, U = 32768$ . Operation Mix 25% `add()`, 25% `rem()`, 50% `con()`

Table 6: Random operation mix benchmark, Intel Xeon system,  $p = 80, c = 1000000, f = 1000, U = 10000$ . Operation Mix 10% **add()**, 10% **rem()**, 80 % **con()**. Operation breakdown: “adds” is the number of successful **add()** operations, “rems” the number of successful **rem()** operations, “cons” the number of element traversals over all **con()** operations, “trav” the number of list element traversals in the search function, “fail” the number of **CAS()** failures, and “rtry” the number of retries in the search function. Variants: a) draconic, b) singly, c) doubly, d) singly-cursor, e) singly-fetch-or, f) doubly-cursor.

Variant	Time (ms)	Total ops	Throughput (Kops/s)	adds	rems	cons	trav	fail	rtry
a)	79940.23	80000000	1000.75	3997685	4002627	159819683981	40004806528	29670	24868
b)	78877.44	80000000	1014.23	3997774	4002741	159771141405	33291359815	24771	1
c)	78152.47	80000000	1023.64	3997991	4002934	159882542042	33279948727	25699	0
d)	56889.93	80000000	1406.22	3994914	3999857	106554989656	26647195005	28602	2
e)	56785.59	80000000	1408.81	3996829	4001795	106534595059	26642102852	28528	1
f)	43498.16	80000000	1839.16	3996761	4001799	106589641350	26636431208	16761	1

Table 7: Deterministic benchmark  $k(i) = i$ , SPARC-T5 system,  $p = 64, n = 100000$ . Operation breakdown: “adds” is the number of successful **add()** operations, “rems” the number of successful **rem()** operations, “cons” the number of element traversals over all **con()** operations, “trav” the number of list element traversals in the search function, “fail” the number of **CAS()** failures, and “rtry” the number of retries in the search function. Variants: a) draconic, b) singly, c) doubly, d) singly-cursor, f) doubly-cursor.

Variant	Time (ms)	Total ops	Throughput (Kops/s)	adds	rems	cons	trav	fail	rtry
a)	514603.24	57600000	111.93	113010	113010	1264254674628	1264624767976	8356	7233
b)	434627.41	57600000	132.53	108395	108395	1275648845411	955962198414	61709	26
c)	407405.79	57600000	141.38	129056	129056	1273179327218	858421075762	168084	0
d)	286471.04	57600000	201.07	108319	108319	38292797	1272875360528	103677	165
f)	173.51	57600000	331962.89	169502	169502	76406418	55414912	719558	4338

Table 8: Deterministic benchmark  $k(i) = t + ip$ , SPARC-T5 system,  $p = 64, n = 10000$ . Operation breakdown: “adds” is the number of successful **add()** operations, “rems” the number of successful **rem()** operations, “cons” the number of element traversals over all **con()** operations, “trav” the number of list element traversals in the search function, “fail” the number of **CAS()** failures, and “rtry” the number of retries in the search function. Variants: a) draconic, b) singly, c) doubly, d) singly-cursor, f) doubly-cursor.

Variant	Time (ms)	Total ops	Throughput (Kops/s)	adds	rems	cons	trav	fail	rtry
a)	4207815.61	5760000	1.37	640000	640000	818179538925	818586650089	4196	3926
b)	3686586.39	5760000	1.56	640000	640000	818018742600	613614553897	33093	12
c)	3687992.32	5760000	1.56	640000	640000	818285400314	613707337468	319493	59
d)	2661259.96	5760000	2.16	640000	640000	204614565262	818219910239	425616	42
f)	968.42	5760000	5947.82	640000	640000	47807228	5779948	288786	62911



Table 9: Random operation mix benchmark, SPARC-T5 system,  $p = 64, c = 1000000, f = 1000, U = 10000$ . Operation Mix 10% `add()`, 10% `rem()`, 80 % `con()`. Operation breakdown: “adds” is the number of successful `add()` operations, “rems” the number of successful `rem()` operations, “cons” the number of element traversals over all `con()` operations, “trav” the number of list element traversals in the search function, “fail” the number of `CAS()` failures, and “rtry” the number of retries in the search function. Variants: a) draconic, b) singly, c) doubly, d) singly-cursor, f) doubly-cursor.

Variant	Time (ms)	Total ops	Throughput (Kops/s)	adds	rems	cons	trav	fail	rtry
a)	40806.24	64000000	1568.39	3200918	3205988	120278045492	30185936395	48740	48521
b)	39319.52	64000000	1627.69	3203939	3208845	120295394002	25351649568	42377	0
c)	40012.58	64000000	1599.50	3200439	3205376	120329461579	25310904222	38018	0
d)	27749.98	64000000	2306.31	3200810	3205729	81939214206	20509159845	58033	2
f)	27130.05	64000000	2359.01	3199878	3204457	87100467429	21785440781	27313	0




Figure 2: Scalability with threads of random operation mix benchmark, Intel Xeon system,  $c = 50000, f = 16384, U = 32768$ . Operation Mix 25% `add()`, 25% `rem()`, 50% `con()`




Figure 3: Scalability with threads of random operation mix benchmark, SPARC-T5 system,  $c = 50000, f = 16384, U = 32768$ . Operation Mix 25% `add()`, 25% `rem()`, 50% `con()`

On average, the speedup in the scalability experiments is about 3.4%. Making the list doubly linked induces additional overhead that actually costs performance. In the scalability experiments, the average speedup of the doubly linked list in comparison to draconic is only about 1.9%.

However, the situation changes when we keep a cursor to the last used item. In case of the doubly linked list this allows us to starting search in the backwards/forwards direction from last item position; for the singly linked list it still cuts the list into two pieces, allowing us to start searching either from the last position or from head, depending on the key of the item to be located. This can improve performance dramatically, in particular for the doubly linked list in the deterministic benchmark (orders of magnitude).

But also in the random mix benchmark the cursor based implementations are significantly faster. When using a cursor, the ability to search backwards in doubly linked list pays off, resulting in significantly better performance, at least on Intel and AMD; on SPARC, the cursor based singly/doubly linked list implementations perform more or less on par.

The mild improvements (with cursor) are easy, unintrusive improvements to the standard, textbook implementation of the lock-free, ordered linked list with significant enough performance improvements to be considered, also for more complex algorithms (skip lists and hash tables) that build on the linked list data structure. These improvements do not comprise the chosen memory reclamation scheme. The approximate backward pointers in the doubly linked improvements can be extremely beneficial in certain cases (the deterministic benchmarks that were designed to highlight this), but come at cost, and complicate memory reclamation. A possible application is a simplification in the implementation of the Stamp-It memory reclamation system [9, 10].

## References

- [1] Tudor David, Rachid Guerraoui, and Vasileios Trigonakis. Asynchronized concurrency: The secret to scaling concurrent search data structures. In *20th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, pages 631–644, 2015.
- [2] Mikhail Fomitchev and Eric Ruppert. Lock-free linked lists and skip lists. In *Proceedings of the Twenty-Third Annual ACM Symposium on Principles of Distributed Computing (PODC)*, pages 50–59, 2004.
- [3] Vincent Gramoli. More than you ever wanted to know about synchronization: synchrobench, measuring the impact of the synchronization on concurrent algorithms. In *20th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming (PPoPP)*, pages 1–10, 2015.
- [4] Rachid Guerraoui and Vasileios Trigonakis. Optimistic concurrency with OPTIK. In *21st ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming (PPoPP)*, pages 8:1–18:12, 2016.
- [5] Timothy L. Harris. A pragmatic implementation of non-blocking linked-lists. In *15th International Conference on Distributed Computing (DISC)*, pages 300–314, 2001.
- [6] Maurice Herlihy and Nir Shavit. *The Art of Multiprocessor Programming*. Morgan Kaufmann Publishers, revised 1st edition, 2012.

- [7] Max Khizhinsky. Cds c++ library, 2020. Retrieved from <https://github.com/khizmax/libcds>.
- [8] Maged M. Michael. High performance dynamic lock-free hash tables and list-based sets. In *Proceedings of the Fourteenth Annual ACM Symposium on Parallel Algorithms and Architectures (SPAA)*, pages 73–82, 2002.
- [9] Manuel Pöter and Jesper Larsson Träff. Poster: Stamp-it, amortized constant-time memory reclamation in comparison to five other schemes. In *23rd ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming (PPoPP)*, pages 413–414, 2018.
- [10] Manuel Pöter and Jesper Larsson Träff. Stamp-it: A more thread-efficient, concurrent memory reclamation scheme in the C++ memory model. In *30th ACM Symposium on Parallelism in Algorithms and Architectures (SPAA)*, pages 355–358, 2018.
- [11] William Pugh. Skip lists: A probabilistic alternative to balanced trees. *Communications of the ACM*, 33(6):668–676, 1990.
- [12] Michael L. Scott. *Shared-Memory Synnchronization*. Synthesis Lectures on Computer Architecture. Morgan & Claypool Publisher, 2013.
- [13] Ori Shalev and Nir Shavit. Split-ordered lists: Lock-free extensible hash tables. *Journal of the ACM*, 53(3):379–405, 2006.
- [14] Håkan Sundell and Philippos Tsigas. Fast and lock-free concurrent priority queues for multi-thread systems. *Journal of Parallel and Distributed Computing*, 65(5):609–627, 2005.

Listing 1: The search function that links out and “physically removes” a marked item, textbook implementation and mild improvements. LOAD and CAS abbreviate the C atomic operations.

```
void pos(long key, list_t *list) {
    node_t *pred, *succ, *curr, *next;

retry:
#ifdef TEXTBOOK
    pred = list->head;
#else
    pred = list->pred;
    if (ismarked(LOAD(&pred->next)) || key <= pred->key)
        pred = list->head;
#endif

    curr = getpointer(LOAD(&pred->next));

    do {
        succ = LOAD(&curr->next);
        while (ismarked(succ)) {
            succ = getpointer(succ);
            if (!CAS(&pred->next, &curr, succ)) {
#ifdef TEXTBOOK
                goto retry;
#else
                next = LOAD(&pred->next);
                if (ismarked(next)) goto retry;
                succ = next;
#endif
            }

            curr = getpointer(succ);
            succ = LOAD(&succ->next);
        }

        if (key <= curr->key) {
            list->pred = pred;
            list->curr = curr;
            return;
        }
        pred = curr;
        curr = getpointer(LOAD(&curr->next));
    } while (1);
}
```

Listing 2: The `rem()` operation, textbook implementation and mild improvement. Also the possible improvement by using an atomic fetch-and-or operation (macro `FAO`) instead of the `CAS`-loop is shown.

```
int rem(long key, list_t *list) {
    node_t *pred, *succ, *node;
    node_t *markedsucc;

    do {
        pos(key, list);
        pred = list->pred;
        node = list->curr;
        if (node->key != key) return 0; // not there

#ifdef TEXTBOOK
        succ = getpointer(LOAD(&node->next)); // unmarked
        markedsucc = setmark(succ);

        if (!CAS(&node->next, &succ, markedsucc))
            continue;
#else
#ifdef FETCH
        succ = FAO(&node->next, MARK_BIT);
        if (ismarked(succ)) return 0;
#else
        succ = LOAD(&node->next);
        do {
            if (ismarked(succ)) return 0;
            markedsucc = setmark(succ);
            if (CAS(&node->next, &succ, markedsucc)) break;
        } while (1);
#endif
#endif

        CAS(&pred->next, &node, succ);
        return 1;
    } while (1);
}
```

Listing 3: The complete search operation with backward pointers. With this search function, `add()` and `rem()` implementations can be kept as they are (textbook). `LOAD`, `STORE` and `CAS` abbreviate the C atomic operations.

```
void pos(long key, list_t *list) {
    node_t *pred, *succ, *curr, *next;

    pred = list->pred;

retry:
    while (ismarked(LOAD(&pred->next)) || key <= pred->key)
        pred = LOAD(&pred->prev);

    curr = getpointer(LOAD(&pred->next));
    do {
        succ = LOAD(&curr->next);
        while (ismarked(succ)) {
            succ = getpointer(succ);
            if (!CAS(&pred->next, &curr, succ)) {
                next = LOAD(&pred->next);
                if (ismarked(next)) goto retry;
                succ = next;
            } else STORE(&succ->prev, pred);

            curr = getpointer(succ);
            succ = LOAD(&succ->next);
        }

        if (LOAD(&curr->prev) != pred)
            STORE(&curr->prev, pred);

        if (key <= curr->key) {
            list->pred = pred;
            list->curr = curr;
            return;
        }
        pred = curr;
        curr = getpointer(LOAD(&curr->next));
    } while (1);
}
```