

Storage Reliability Calculations

02 December 2016

When designing storage systems, for example comparing different coding schema, it is necessary to calculate reliability numbers. Here I summarize the useful concepts from reliability engineering and their math relations.

Reliability engineering

```
 $\lambda$  = failure rate, expressed in failures per unit of time  
MTBF (mean time between failures) =  $1/\lambda$  // why? see hazard function and exponential distribution  
 $AFR = 1 - e^{(-exposure\_time/MTBF)} = 1 - e^{(-exposure\_time*\lambda)}$  = probability of failure occurs within the exposure_  
Assuming a small AFR (<5%), we can approximatedly have  $AFR \approx 8760/MTBF$   
 $1 - AFR$  = probability of failure does not occur (survived) within the exposure_time
```

Hazard function

T = a random variable of the time until some event of interest happens
 $f(t)$ = probability density of event happens at t
 $F(t) = P(T \leq t)$ = probability of event happens before t

$S(t)$ = survivor function / reliability function := $1 - F(t)$ = probability of event happens after t , or say even

$h(t)$ = hazard function := $\lim_{h \rightarrow 0^+} P[t \leq T < t + h | T \geq t] / h = f(t) / S(t) = -d \ln(1 - F(t)) / dt = -d \ln(S(t))$
 $h(t)dt = f(t)dt / S(t) \approx P[\text{fail in } [t, t+dt) \mid \text{survive until } t]$
 $H(t)$ = cumulative hazard function := $\int (0 \rightarrow t, h(u)du) (t > 0) = -\ln(1 - F(t)) = -\ln(S(t))$
 So, $S(t) = e^{-H(t)}$, $f(t) = h(t)e^{-H(t)}$

Exponential Distribution: $T \sim \text{Exp}(\lambda)$, for $t > 0$.
 $f(t) = \lambda e^{-\lambda t}$ for $\lambda > 0$ (scale parameter)
 $F(t) = 1 - e^{-\lambda t}$ // this is the AFR mentioned above
 $S(t) = e^{-\lambda t}$ // this is the (1 - AFR) mentioned above
 $h(t) = \lambda$ // constant hazard function. this is the failure rate mentioned above
 $H(t) = \lambda t$
 Characteristics:
 $E(T) = 1/\lambda$ // this is the MTBF mentioned above
 "Lack of Memory": $P[T > t] = P[T > t + t_0 | T > t_0]$. Probability of surviving another t time units does not depend on how long you have already survived.

Reference: [1] (<http://web.stanford.edu/~lutian/coursepdf/unit1.pdf>)

Calculating MTBF from observation

observed AFR = failed device count / total device count
 calculated MTBF(hrs) = 8760 / AFR // 8760 is the number of hours in a year
 exposure_time: for storage system with auto repair, the exposure time of second node failure should be the recovery
 failure possibility in a given time window = $F(t) = 1 - e^{-\lambda t} \approx \lambda t = t / \text{MTBF} = t * \text{AFR} / 8760$, t is the time window

Availability vs durability

Availability: the probability of data is accessible by user. The data needs to be on disk, and server nodes have
 Durability: the probability of data is safe on disk. The data may not be available to user, but still on disk.
 For example: AWS S3 offers 99.999999999% durability and 99.9% availability.

Data loss probabilities

Basics

$A(n, k) := \text{Permutation}(n, k) = n! / (n-k)!$
 $C(n, k) := \text{Combination}(n, k) = n! / ((n-k)! k!)$
 $Z_m := \{0, 1, \dots, m-1\}$
 $\text{Subset}(A) := \{a \mid a \text{ is subset of } A\}$
 $\neg A := \text{overbar } A := \text{complementary set of } A$
 $I := \text{the complete set} = A + \neg A$
Code schema (D-K, K) := the simplest erasure-coding schema of D-K data fragments and K parity fragments. Wh
P(event): a probability is operating on event. An event is a set. So set algebra applies to events.
n: event1 AND. This is aligned with set algebra
U: event1 OR. This is aligned with set algebra
 \neg or overbar: event NOT happen. This is aligned with set algebra

Basic formulas (proof omitted)

Event breakdown formula: (aux1)
Multiple event breakdown formula: (aux2)
If A_i are mutually independent, then $\neg A_i$ are mutually independent: (aux3)
Conditional probability can reverse: (aux4)
If A_i are mutually conditional independent given N, then $\neg A_i$ are mutually conditional independent given N:
Space division, which is useful for node failure probability: (aux6)

Notations

Cluster setup:

We have N nodes, each have independent probability p to fail.
We have M extents (extent is the basic data unit, i.e. object), each is encoded into a D-fragment code.
We randomly place the fragments on nodes, each node holds at most 1 fragment from a given extent.
The code schema can be complex, e.g. LRC code.
The placement may have constraints, e.g. domains, copysets.

Node failure

Node set is Z_N ("N" in subscript)
 $NL :=$ a node failure event. NL is a subset of Z_N .
 $NLS :=$ the space of NL = $\text{Subset}(Z_N)$

Extent and Extent Loss

$E_i :=$ Extent i, i in Z_M ("M" in subscript)
 $EL :=$ an extent loss event = {indexes of the lost fragments} // fragment index is element of Z_D ("D"
 $ELS :=$ the space of EL = {EL|any EL}, it is the set of which fragment-loss can cause extent loss
 $\neg ELS := \text{Subset}(Z_D) - ELS$, it is the set of safe fragment-loss

Placement

$PM :=$ a placement = [node indexes of the placed fragment, for fragment from $idx=0$ to $idx=D-1$]
 $PM_i :=$ extent i's placement

S_{all} ("all" in subscript) := {all possible placements }

S := set of the allowed placements, S is subset of S_{all} , due to placement constraints

SNL ("NL" in subscript) := {placement that will cause extent loss}, given NL

nh ("h" in subscript) := set intersect and map node index to fragment index. if $(PM \cap nh \cap NL)$ is element of

Data loss

$E_i \cap DL$:= an event of extent i is lost. $P(E_i \cap DL)$ is the probability of lost extent i

$E_i \cap DL \mid NL$ = extent i loses, given NL. $P(E_i \cap DL \mid NL)$ is the conditional probability of $E_i \cap DL$ given NL

DL := an event of data loss. $P(DL) = P(DL > 0)$ is the probability of losing at least 1 extent in the cluster

Assumptions

Extent placements are mutually independent: (asp1)

Given an extent, placements are exclusive: (asp2)

Placement is evenly randomized: (asp3)

Node failure is exclusive: (asp4)

Node failure is independent with extent placement: (asp5)

Implications (proof omitted)

Extent placement sets are independent: (imp1)

Node failure is independent with extent placement sets: (imp2)

Sample space (elementary event space)

An elementary event is: $(NL, PM_0, PM_1, \dots, PM_{M-1})$ ("0", "1", ..., "M-1" in subscript)

The sample space is: $NLS \times S \times S \times \dots \times S$ (count of $S = M$)

Given NL, $P(\emptyset)$:= elementary event probability, $P(\emptyset)$ formula: (imp3)

Key takeaway: most of the proofs below is conducted by breaking events down to elementary events
probability is essentially counting elementary events

Conclusions (proof omitted)

Calculate $P(E_i \cap DL \mid NL)$: (thr1)

$(E_i \cap DL \mid NL)$ are mutually conditional independent: (thr2)

$(E_i \cap DL \mid NL)$ and $\neg(E_i \cap DL \mid NL)$ are mutually conditional independent: (thr3)

Calculate $P(E_i \cap DL)$: (thr4) // To prove the two forms are equal, breakdown them into $P(\emptyset)$

$E_i \cap DL$ are not independent: (thr5)

Alternative assumption, $E_i \cap DL$ are still not independent: Assume extent placements are fixed/known, the sample

Data loss count distribution: (thr6)

Data loss probability: (thr7)

$P(E_i \cap DL)$ vs $P(DL)$: connected by $|SNL|/|S|$ ("NL" in subscript), but not always matched

Counter-case: from $(N, p, D, K, M) = (20, 0.16, 4, 1, 50)$ to $(20, 0.16, 12, 2, 50)$, $P(DL)$ drops, but $P(E_i \cap DL)$ vs $P(DL)$: meaning

$P(E_i \cap DL)$: when a customer asks how easy is his extent to be lost.

$P(DL)$: when the operation asks how easy will they encounter data loss alert, and how much lost data.

Questions

If we have a large number of nodes, and many extents, will the extent loss event become nearly independent?

We can think this way

Since cluster is large, $(E_i \text{ DL})$ is less affected by $(E_j \text{ DL})$.

When we take extents E_0, E_1, \dots, E_{m-1} , they are more likely to be not overlapped

When E_0, E_1, \dots, E_{m-1} are not overlapped, we can deduce that $E_i \text{ DL}$ are independent

Below are the formulas stated above

(aux1) Given $B_i \cap B_j = \emptyset, \forall i \neq j \in Z_m$, and $\bigcup_{i \in Z_m} B_i = I$

$$P(A) = \sum_{i \in Z_m} P(B_i \cap A) = \sum_{i \in Z_m} P(B_i)P(A|B_i)$$

(aux2) Given $A_{i,j}, i \in Z_n, j \in Z_{m_i}; A_{i,j_1} \cap A_{i,j_2} = \emptyset, \forall i \in Z_n, \forall j_1 \neq j_2 \in Z_{m_i}$. We have

$$P(\bigcap_{i \in Z_n} \bigcup_{j \in Z_{m_i}} A_{i,j}) = \sum_{j_0 \in Z_{m_0}} \sum_{j_1 \in Z_1} \dots \sum_{j_{n-1} \in Z_{m_{n-1}}} P(\bigcap_{i \in Z_n} A_{i,j_i})$$

If A_{i_1,j_1} and A_{i_2,j_2} are independent, $\forall i_1 \neq i_2 \in Z_n, \forall j_1 \in Z_{m_{i_1}}, j_2 \in Z_{m_{i_2}}$. We have

$$P(\bigcap_{i \in Z_n} \bigcup_{j \in Z_{m_i}} A_{i,j}) = \sum_{j_0 \in Z_{m_0}} \sum_{j_1 \in Z_1} \dots \sum_{j_{n-1} \in Z_{m_{n-1}}} \prod_{i \in Z_n} P(A_{i,j_i})$$

(aux3) Given A_0, A_1, \dots, A_{m-1} are mutually independent. We have

$$P(\overline{A_0} \cap \overline{A_1} \cap \dots \cap \overline{A_{m-1}}) = \prod_{i \in Z_m} P(\overline{A_i})$$

(aux4) $P(\overline{A} | N) = 1 - P(A | N)$

(aux5) Given A_0, A_1, \dots, A_{m-1} are mutually conditional independent given N . We have

$$P(\overline{A_0} \cap \overline{A_1} \cap \dots \cap \overline{A_{m-1}} | N) = \prod_{i \in Z_m} P(\overline{A_i} | N)$$

(aux6) Given $A_i \cap A_j = \emptyset, \forall i \neq j \in Z_m, \bigcup_{i \in Z_m} A_i = S, P(e) = p, \forall e \in S$. Event $e \in S$ are mutually independent. We have

$$P(e \in \{a_0 \cup a_1 \cup \dots \cup a_{m-1} | a_0 \in SA_1, a_1 \in SA_2, \dots, a_{m-1} \in SA_{m-1}\}) = \prod_{i \in Z_m} \sum_{a_i \in SA_i} p^{|a_i|} (1-p)^{|SA_i| - |a_i|}$$

(asp1) $P(\bigcap_{i \in J} PM_i) = \prod_{i \in J} P(PM_i), \forall J \subseteq Z_M$

(asp2) $P(PM_i = PM_1 \cap PM_i = PM_2) = 0, \forall PM_1 \neq PM_2 \in S$

(asp3) $P(PM_i) = \frac{1}{|S|}, \forall i \in Z_M$

(asp4) $P(NL = NL_1 \cap NL = NL_2) = 0$

(asp5) $P(NL \cap (\bigcap_{i \in J} PM_i)) = P(NL)P(\bigcap_{i \in J} PM_i), \forall NL \in NLS, J \subseteq Z_M$

(imp1) $P(\bigcap_{i \in J} PM_i \in S_i) = \prod_{i \in J} P(PM_i \in S_i) = \sum_{e_0 \in S_0} \sum_{e_1 \in S_1} \dots \sum_{e_{m-1} \in S_{m-1}} \prod_{i \in Z_m} P(PM_i = e_i) = \left(\frac{1}{|S|}\right)^{|J|} \prod_{i \in J} |S_i|, \forall J \subseteq Z_M$

(imp2) $P(NL \cap (\bigcap_{i \in J} PM_i \in S_i)) = P(NL)P(\bigcap_{i \in J} PM_i \in S_i), \forall J \subseteq Z_M$

(imp3) $P(\odot) = p^{|NLS|} (1-p)^{N-|NLS|} \frac{1}{|S|}$

(thr1) $P(E_i DL | NL) = \frac{|S_{NL}|}{|S|}$

(thr2) $P(\bigcap_{i \in Z_m} E_i DL | NL) = \prod_{i \in Z_m} P(E_i DL | NL)$

(thr3) $P(\bigcap_{0 \leq i \leq r} E_i DL \cap \bigcap_{r+1 \leq i \leq m-1} \overline{E_i DL} | NL) = \prod_{0 \leq i \leq r} P(E_i DL | NL) \prod_{r+1 \leq i \leq m-1} P(\overline{E_i DL} | NL)$

(thr4) $P(E_i DL) = \sum_{NL \in NLS} P(NL) \frac{|S_{NL}|}{|S|} = \sum_{EL \in ELS} p^{|EL|} (1-p)^{D-|EL|}$

(thr5) $P(\bigcap_{i \in Z_m} E_i DL) = \sum_{NL \in NLS} P(NL) \left(\frac{|S_{NL}|}{|S|}\right)^m \neq \left(\sum_{NL \in NLS} P(NL) \frac{|S_{NL}|}{|S|}\right)^m = \prod_{i \in Z_m} P(E_i DL)$

(thr6) $P(DL = r) = C_M^r \sum_{NL \in NLS} P(NL) \left(\frac{|S_{NL}|}{|S|}\right)^r \left(1 - \frac{|S_{NL}|}{|S|}\right)^{M-r}$

(thr7) $P(DL) = 1 - P(DL=0) = 1 - \sum_{NL \in NLS} P(NL) \left(1 - \frac{|S_{NL}|}{|S|}\right)^M$

Reliability Limits

I want to explore what limits the reliability of super large storage clusters.

Basics

$E(a) :=$ expectation of random_variable

$E(a|case) :=$ conditional expectation of random_variable = $\sum(v \text{ in } \{a\text{'s possible values}\}, v * P(a=v|case))$

Basic formulas

Conditional expectation: (aux7)

Notations

$\lambda :=$ node failure rate = $d(P(\text{node fail in time } dt))/dt =$ a constant

$dt :=$ a very small time window, which is usually used with λ

$\lambda dt =$ node failure possibility in time window dt

$K :=$ K that, an extent losing $l > K$ fragments is never recoverable, losing $l \leq K$ fragments can be recoverable
usually, $K =$ parity fragment count

$SR :=$ storage overhead = $D/(D-K)$

$SD :=$ user data size per extent

$TD :=$ total user data size

$DS :=$ fragment size = $SD * SR / D = SD / (D-K)$

Assumptions

Breakdown extents by how many fragments lost: (asp6)

M_l ("l" in subscription) := extent count of losing l fragments

$NB :=$ network bandwidth for data recovery. note that other network bandwidth can be used for customer t

NB_l ("l" in subscription) := network bandwidth for recover extents in M_l

rv_l ("l" in subscription) := given the code schema, in all cases of losing l fragments, ratio of the re

Conclusions

First, we study the the stable state. I.e. NB_l needs to be proper to hold M_l stable.

Extent count that flows into M_i due to node loss: (thr8)

Extent count that flows into M_i due to recovery: (thr9)

Extent count that flows out of M_i due to recovery: (thr10)

Extnet count that flows out of M_i due to node loss: (thr11)

By balancing the flow in and flow out, to maintain the stable state, we have: (thr12). Conclusions

It is not possible to balance. Not possible even $M_0=K$, $M_1..M_K = 0$. There is no balance state.

If we force $M_0..M_{K-1}$ to balance, M_K will drop. The process repeat until all $M_K..M_1$ drop to 0, survi

If we force $M_K..M_1$ to balance, M_0 will drop. The process repeats until all $M_0..M_{K-1}$ drop to 0, surv

A healthy cluster can only be $M_0=M$: (thr13)

The other $M_1..M_K$ should be infinitely small, ignorable count

Next, let's find out data loss rates

Extent loss rate $\lambda(\text{ext})$: (thr14)

Cluster data loss rate $\lambda(\text{DL})$: (thr15)

$\lambda(\text{ext})$ and $\lambda(\text{DL})$ are the solo result of M_i states. They have no relation with NB_i . While NB_i is what
When $M_0=M$, $\lambda(\text{ext})$ and $\lambda(\text{DL})$ are the best of what they can get from NB_i .

Practical markov for single extent loss probability: (thr16)

Notations

Let $NBS_i :=$ the available network recovery bandwidth for this single extent, when it has lost i

The bandwidth is affected by repair priority, TOR and node NIC bandwidth.

$P(\text{EDL } M_i) :=$ probability of lose this extent when it is at M_i

$P(\text{EDL } M_0)$ is what we try to calculate

Note that this formula is not math strict, because in an absorbing markov chain the extent loss probability

https://www.dartmouth.edu/~chance/teaching_aids/books_articles/probability_book/Chapter11.pdf

But we are here to calculate the "practical" probability

This formula assumes we follow the repair lifecycle of the given extent

While the above $\lambda(\text{ext})$ assumes we take any extent from the cluster, which has M_1 - M_K count is in

This formula is useful because

It takes network recovery bandwidth into consideration, very practical

It follows the entire repair lifecycle of the given extent

Conclusions

To maintain cluster healthy, recovery bandwidth NB should be no less than $\lambda \cdot TD \cdot SR$.

The necessary recovery bandwidth grows with total user data size, code schema storage overhead, but
Recovery bandwidth NB limits how much data can be there in a super large storage cluster

λ and code schema storage overhead multiply it by factors

Cluster node count has no relation with NB . But user data size does.

For customer-faced reliability, extent loss rate $\lambda(\text{ext})$ should be low. It is the direct result of λ and code

For operation ease, cluster data loss $\lambda(\text{DL})$ should be low. It is the direct result of λ and code schema

$\lambda(\text{ext})$ and $\lambda(\text{DL})$ have no direct relationship with NB . But NB is necessary to pull extents from M_1 .

If $M_1..M_K \neq 0$, the $\lambda(\text{ext})dt$ and $\lambda(\text{DL})dt$ raise by factor of $1/(\lambda dt)$

Below are the formulas stated above

(aux7) Given $B_i \cap B_j = \emptyset, \forall i \neq j \in Z_m$, and $\bigcup_{i \in Z_m} B_i = I$

$$E(a) = \sum_{v \in \{a's \text{ possible values}\}} v * P(a = v) = \sum_{i \in Z_m} P(B_i) E(a|B_i)$$

(asp6) $M = M_0 + M_1 + \dots + M_l + \dots + M_D; M_l = 0, \forall l > K$

$$NB = NB_0 + NB_1 + \dots + NB_l + \dots + NB_D; NB_0 = 0; NB_l = 0, \forall l > K; \forall l, \text{ if } M_l = 0, \text{ then } NB_l = 0$$

$$rv = [rv_0, rv_1, \dots, rv_l, \dots, rv_D]; rv_0 = 1; rv_l = 0, \forall l > K; rv_{l+1} < rv_l, \forall l$$

$$\begin{aligned} \text{(thr8)} \quad CntExt_LeftFlowIn_i &= rv_i \sum_{|NL|=0}^N \sum_{l=0}^{i-1} (\lambda dt)^{|NL|} (1 - \lambda dt)^{N-|NL|} * M_l * C_{D-l}^{i-l} C_{N-(D-l)}^{|NL|-(i-l)} \\ &\approx rv_i * \lambda dt * M_{i-1} * (D - i + 1) \end{aligned}$$

$$\begin{aligned} \text{(thr9)} \quad CntExt_RightFlowIn_i &= \sum_{l=i+1}^D M_l * \prod_{j=i+1}^l \frac{NB_j}{M_j * DS} * (dt)^{l-i} \\ &\approx \frac{NB_{i+1}}{DS} dt \end{aligned}$$

$$\text{(thr10)} \quad CntExt_LeftFlowOut_i = \frac{NB_i}{DS} dt$$

$$\begin{aligned} \text{(thr11)} \quad CntExt_RightFlowOut_i &= M_i \sum_{|NL|=0}^N (\lambda dt)^{|NL|} (1 - \lambda dt)^{N-|NL|} (C_N^{|NL|} - C_{N-(D-i)}^{|NL|}) \\ &\approx M_i * \lambda dt * (D - i) \end{aligned}$$

$$\text{(thr12)} \quad CntExt_LeftFlowIn_i + CntExt_RightFlowIn_i = CntExt_LeftFlowOut_i + CntExt_RightFlowOut_i, \forall i=0..K$$

$$\Rightarrow \left[\begin{array}{l} \frac{NB_k}{DS} * \frac{1}{\lambda D} = rv_k * M_{k-1} * \left(1 - \frac{k-1}{D}\right) - M_k * \left(1 - \frac{K}{D}\right) \\ \frac{NB_i - NB_{i+1}}{DS} * \frac{1}{\lambda D} = rv_i * M_{i-1} * \left(1 - \frac{i-1}{D}\right) - M_i * \left(1 - \frac{i}{D}\right) \end{array} \right] \Rightarrow$$

$$\text{with } \frac{NB_0}{DS} * \frac{1}{\lambda D} = 0 \quad \text{and} \quad \forall l, \text{ if } M_l = 0, \text{ then } NB_l = 0$$

$$\text{Let } b_i = \frac{NB_i}{DS} * \frac{1}{\lambda D}, a_i = M_i * \left(1 - \frac{i}{D}\right)$$

$$\Rightarrow \left[\begin{array}{l} b_k = rv_k a_{k-1} - a_k \\ b_i - b_{i+1} = rv_i a_{i-1} - a_i \\ b_0 = 0 \\ \forall l, \text{ if } a_l = 0 \text{ then } b_l = 0 \end{array} \right] \Rightarrow$$

If start from $b_k = rv_k a_{k-1} - a_k$, we get $b_{i, \text{from } b_k} = rv_i a_{i-1} - \sum_{l=i+1}^K (1 - rv_l) a_{l-1} - a_k$, then

$$b_{0, \text{from } b_k} - b_0 = -\left(\sum_{l=1}^K (1 - rv_l) a_{l-1} + a_k\right) \leq 0$$

If start from $b_0 = 0$, we get $b_{i, \text{from } b_0} = \sum_{l=1}^{i-1} (1 - rv_l) a_{l-1} + a_{i-1}$, then

$$b_{k, \text{from } b_0} - b_k = \sum_{l=1}^K (1 - rv_l) a_{l-1} + a_k \geq 0$$

If we balance the equation group, then we get

$$M_l = 0, NB_l = 0, \forall l = 0..K \text{ (note that } l=0 \text{ is included)}$$

$$\text{(thr13)} \quad M_0 = M$$

$$\frac{NB}{DS} dt = M_0 * \lambda dt * D \Rightarrow NB = \lambda * D * DS * M = \lambda * M * SD * SR = \lambda * TD * SR$$

$$(thr14) \lambda_{ext} dt = \sum_{i=0}^D \frac{M_i}{M} \sum_{l=1}^{D-i} C_{D-i}^l (\lambda dt)^l (1 - \lambda dt)^{(D-i)-l} (1 - rv_{i+l}) \approx \lambda D * dt * \sum_{i=0}^D \frac{M_i}{M} (1 - \frac{i}{D}) (1 - rv_{i+1})$$

$$\lambda_{ext} dt = \sum_{l=1}^D C_D^l (\lambda dt)^l (1 - \lambda dt)^{D-l} (1 - rv_l), \text{ if } M_0 = M$$

$$(thr15) \lambda_{DL} dt \approx 1 - \sum_{|NL|=0}^N C_N^{|NL|} (\lambda dt)^{|NL|} (1 - \lambda dt)^{N-|NL|} \prod_{i=0}^D (\sum_{l=1}^{D-i} \frac{C_{D-i}^l C_{N-(D-i)}^{|NL|-l}}{C_N^{|NL|}} rv_{i+l} + \frac{C_{D-i}^0 C_{N-(D-i)}^{|NL|}}{C_N^{|NL|}}) M_i$$

// The above uses “≈” because I averaged Π extent loss probability by rv

$$\approx \lambda * N * dt * (1 - \prod_{i=0}^D (1 - \frac{D-i}{N} (1 - rv_{i+1})))^{M_i}$$

$$\lambda_{DL} dt \approx 1 - \sum_{|NL|=0}^N C_N^{|NL|} (\lambda dt)^{|NL|} (1 - \lambda dt)^{N-|NL|} \left(\sum_{l=0}^D \frac{C_D^l C_{N-D}^{|NL|-l}}{C_N^{|NL|}} rv_l \right)^M, \text{ if } M_0 = M$$

$$(thr16) P(EDL M_i) = \sum_{l=i+1}^D C_{D-i}^{l-i} \left(\lambda \frac{NBS_i}{DS} \right)^{l-i} \left(1 - \lambda \frac{NBS_i}{DS} \right)^{D-l} ((1 - rv_l) + rv_l P(EDL M_l))$$

$$P(EDL M_K) = 1 - \left(1 - \lambda \frac{NBS_K}{DS} \right)^{D-K}$$

Reliability framework

To summarize the technique framework to ensure storage reliability

- * Data input: ensure the reliability when data is coming in
 - * End-to-end verification: ensure what user writes is what is stored
 - * Commit safe: when user is ack-ed, ensure the data is guaranteed persistent
- * Data storing: ensure the reliability at where data is stored
 - * Good data should not be overwritten. In case a bug or at least we can recover data from history
 - * Data recovery
 - * Detect data loss as fast as you can
 - * On-the-fly, when user read data, from node loss, from disk failure, etc. Summarize whatever info
 - * Prediction, from unstable/unhealthy/untrustable nodes that are going to loss, from disks that fail
 - * Repair action should be prioritized
 - * Repair the data at high risk first.
 - * Repair actions should be well-scheduled and well-managed.
 - * The overall reliability is determined by how fast you can repair.
 - * Watch out that your data is more vulnerable when server upgrading, because
 - * Some portion of data is unavailable
 - * The user traffic becomes larger because of reconstruction read
 - * The repair traffic also becomes larger
 - * Code schema (e.g. 3-replica, erasure-coding, LRC) is very important for data reliability.
 - * It is also important for balancing reliability with overhead of storage, bandwidth, latency, etc.
 - * Policed auto transition between different code schemas may be necessary.
- * Silent data corruption: This is a real problem. Background check and repair can help; but too time/resource consuming
 - * Some piece hardware from vendors can have problem. Some firmware can have problem, go wrong, or have bugs
 - * Old clusters are highly risked: old devices are easy to fail, and in burst, which generates a lot of data corruption
- * Data serving
 - * Data availability and data durability are different. Good data serving ensures data availability
 - * Nodes can be transiently unavailable. Detect transient and permanent before you head to expensive data
 - * Node failure can usually be repaired quickly by a reboot. Node can jump online and offline. Be careful
 - * Customer impact is different, i.e. even data is lost, if we are in the lucky time window where customer is not using it
- * Raw disk recovery: when everything is gone, we expect to extract user data from a raw disk
 - * Make sure you have a way to know the raw binary format on disk. Extract data from raw disk can be the first step
 - * Make sure the necessary metadata for raw disk recovery can be obtained and well protected.
- * Hardware: hardware failure rate has direct impact on everything about reliability, very straightforward. Be aware of it

Other references

- DataDomain DIA (<http://www.emc.com/collateral/software/white-papers/h7219-data-domain-data-invol-arch-wp.pdf>)

storage²² (/categories.html#storage-ref)

storage³² (/tags.html#storage-ref)

reliability⁷ (/tags.html#reliability-ref)

cloud¹⁵ (/tags.html#cloud-ref)

[← Previous \(/storage/Tail-Latency-Study\)](/storage/Tail-Latency-Study)

[Archive \(/archive.html\)](/archive.html)

[Next → \(/storage/Reading-Fast16-Papers\)](/storage/Reading-Fast16-Papers)

Create an Issue () or comment below

© 2022 Accela Zhao (/about.html) All rights reserved.